

Contents

CTF Crypto Writeups: fence, mario, linchan, hackel	1
1. fence — NTRU-style lattice key recovery	1
2. mario — UOV (Unbalanced Oil-and-Vinegar) key recovery	4
3. linchan — hidden matrix-conjugacy (MinRank + Sylvester-equation key recovery)	6
Common thread across all three	9
4. hackel — semidirect-product group-presentation key recovery	9
Common thread across all four	13
Appendix — Screenshots from the challenge session	13

CTF Crypto Writeups: fence, mario, linchan, hackel

Four independent challenges, each built around a different piece of “post-quantum-flavored” or “modern-algebra-flavored” mathematics hiding a classic implementation weakness. The first three were solved end-to-end from source + output files; the fourth (hackel) is reconstructed from a static session transcript only (no live server access, no source code) — see §4 for what is independently verified vs. reasoned from structure.

The first three were solved by:

1. Reading the generator script carefully to understand exactly what algebraic object is being hidden and how the output is encoded.
2. Identifying an invariant or leak that lets us recover the hidden object without brute forcing the whole secret space.
3. Writing Python (with `fpylll` for lattice reduction where needed) to carry out the recovery, then re-deriving the symmetric key and decrypting.

1. fence — NTRU-style lattice key recovery

1.1 Files

- `fence.py` — the generator/encryption script.
- `flag.enc` — its JSON output (5 independent “rounds” of public data + ciphertext).

1.2 Understanding the scheme

`fence.py` works in the ring $R_q = \mathbb{Z}_q[x]/(x^{n+1})$ with:

```
n = 128
q = 268435361      # prime, ~2^28
w = 80             # Hamming weight of secret ternary polynomials
r = 5              # number of independent "rounds"
```

Key helper functions:

- `pm(a, b)` — **negacyclic polynomial multiplication** mod x^{n+1} (this is exactly the NTRU ring product: $x^n \equiv -1$).
- `iv(a)` — polynomial **inversion** mod (q, x^{n+1}) via an extended-Euclidean algorithm over $\mathbb{Z}_q[x]$.
- `gn()` — generates a random **ternary** polynomial of length $n=128$ with exactly $w/2$ coefficients $+1$, $w/2$ coefficients -1 , and the rest 0 (Hamming weight $w=80$).

- $ky(a, b, s)$ — derives a symmetric key from the **private** polynomials a, b and a random salt s (note: this is *not* a public-key decryption step; whoever calls ky must already know a and b).
- en/dc — AES-GCM-like authenticated stream cipher (SHAKE-256 keystream + HMAC tag) keyed by the output of ky .

`main()` does, for idx in $range(r=5)$:

```
a, b = gn(), gn()
h = pm(b, iv(a))          # PUBLIC value:  $h = b * a^{-1} \pmod{q, x^{n+1}}$ 
m = <mask or (mask  $\oplus$  flag)> # see below
hs.append(h)
cs.append(en(a, b, h, m))  # ciphertext, keyed via  $ky(a, b, s)$ 
```

The **secret sharing** structure:

```
acc = [0]*len(flag)
for idx in range(r):
    ...
    if idx < r - 1:
        m = random_bytes          # a fresh one-time pad
        acc ^= m
    else:
        m = acc ^ flag            # last round: pad XORed with the flag
```

So $flag = acc \oplus m_{\{r-1\}}$ where $acc = m_0 \oplus m_1 \oplus m_2 \oplus m_3$ (the first four rounds' random pads). **To recover the flag we must decrypt all 5 rounds**, not just the last one — each round's plaintext m_{idx} is needed.

Each round publishes $h = b \cdot a^{-1} \pmod{q, x^{n+1}}$ — this is precisely the **NTRU public key** construction, with a playing the role of the “denominator”/f polynomial and b the “numerator”/g polynomial. a, b are never published; only h .

Since en/dc derive their key directly from the private a, b (via ky), decrypting any round's ciphertext requires recovering that round's private key pair (a, b) from its public h . **This is the classical NTRU private-key recovery problem** — a lattice problem.

1.3 The lattice attack

For small n (here $n = 128$), the NTRU secret polynomials are so much shorter than a “generic” lattice vector that even plain **LLL / light BKZ** lattice reduction recovers them directly.

Step 1 — build the NTRU lattice. For a public h , define the negacyclic shift matrix M (an $n \times n$ matrix over $\mathbb{Z}$) where row i holds the coefficients of $x^i \cdot h(x) \pmod{x^n + 1}$, reduced into **balanced representatives** $[-q/2, q/2)$ (this balancing turns out to matter a lot for reduction quality). Then build the standard $2n \times 2n$ NTRU lattice basis:

$$\begin{bmatrix} I_n & M \\ 0 & q \cdot I_n \end{bmatrix}$$

Any vector $(a \parallel b)$ with $a \cdot h \equiv b \pmod{q, x^{n+1}}$ is a lattice point (choose the correct integer multiples of q to cancel the reduction). Since $h = b \cdot a^{-1}$, our secret (a, b) — with entries in $\{-1, 0, 1\}$ and only $w = 80$ nonzero coordinates each, Euclidean norm $\approx \sqrt{160} \approx 12.6$ — is **enormously shorter** than a “random” vector in this 256-dimensional lattice, whose Gaussian-heuristic expected shortest-vector length is:

$$\lambda_1 \approx \sqrt{(2n / 2\pi e) \cdot \det(L)^{1/2n}} = \sqrt{(256/17.08) \cdot \sqrt{q}} \approx 3.87 \times 16384 \approx 63,000$$

A gap of $\sim 5000\times$ between the target and the “generic” shortest vector means the private key is trivially the shortest vector once the lattice is reduced.

Step 2 — reduce with fpylll.

```
from fpylll import IntegerMatrix, LLL, BKZ

def negacyclic_shift(h, i):
    # returns balanced-representative coefficients of  $x^i * h(x) \bmod (x^N + 1)$ 
    ...

def build_lattice(h):
    M = IntegerMatrix(2*N, 2*N)
    for i in range(N):
        M[i, i] = 1
        row = negacyclic_shift(h, i)
        for j in range(N):
            M[i, N+j] = row[j]
    for i in range(N):
        M[N+i, N+i] = Q
    return M

def recover(h):
    M = build_lattice(h)
    LLL.reduction(M)
    best = shortest_row(M)
    for bs in (10, 20, 30, 40, 50, 60):      # escalate BKZ block size if needed
        if best_norm2 <= target:
            break
        BKZ.reduction(M, BKZ.Param(block_size=bs))
        best = shortest_row(M)
    return best
```

Plain **LLL alone landed around $\text{norm}^2 \approx 15,000\text{--}23,000$** (rows of $\sim 120\text{--}150$ in length) — better than random but not yet the actual key. Running a light **BKZ with block size 10** on top pushed every one of the 5 instances straight down to **$\text{norm}^2 = 160$** — exactly $80 + 80$ — confirming exact recovery of the weight-80 ternary (a, b) pair, for **all 5 rounds**.

Step 3 — verify and decrypt.

For each recovered (a, b) :

```
assert pm(b, iv(a)) == h      # confirms this is really the private key for this h
m = dc(a, b, h, ciphertext)    # decrypts this round's plaintext
```

Step 4 — reassemble the flag.

```
acc = bytearray(len(m))
for idx in range(5):
    m = decrypt(idx)
    if idx < 4:
        acc ^= m          # accumulate the one-time pads
    else:
        flag = acc ^ m    # final round unmask the flag
```

1.4 Result

ASIS{qu4ntum_c0h3r3nc3_1n_0v3r5tr3tch3d_h4rm0n1c_f13ld5!}

1.5 Why it broke

$n = 128$ is far too small for NTRU-style lattice hardness assumptions to hold against LLL/BKZ — real NTRU parameter sets use $n \geq 500+$. The 5-way secret-sharing across rounds didn't add real security either, since **every** round's h is independently, trivially breakable, so all the pieces needed to reconstruct the flag are recoverable.

2. mario — UOV (Unbalanced Oil-and-Vinegar) key recovery

2.1 Files

- mario.py — generator script.
- output.txt — JSON output (public quadratic system + leaked “reports” + AES-GCM ciphertext).

2.2 Understanding the scheme

This is a **multivariate quadratic (MQ) signature-style scheme** over $\text{GF}(16)$ (elements 0–15, represented as nibbles; addition is XOR, multiplication via a precomputed table MUL using an explicit modulus polynomial x^4+x+1).

Parameters:

```
n, m, d, s = 96, 72, 24, 64
v = n - d = 72      # "vinegar" variables
```

Secret key: k_{mat} , a $v \times d$ matrix over $\text{GF}(16)$, defining a linear embedding

```
def oil_embed(k_mat, x):      # x has d=24 coordinates
    y[:v] = k_mat @ x         # "vinegar" part is a linear function of x
    y[v:] = x                 # "oil" part is x itself
    return y                  # y has n=96 coordinates
```

The image of `oil_embed` (as x ranges over $\text{GF}(16)^{24}$) is a **24-dimensional linear subspace 0 of $\text{GF}(16)^{96}$** — the classic UOV **oil subspace**, just “twisted” so it’s not simply the last 24 coordinates.

Public key construction (`build_public_map`): for each of $m = 72$ quadratic forms, random coefficients are assigned to every term touching a vinegar variable ($i < v$ or $j < v$), and then the **oil-oil coefficients are back-solved** so that the quadratic form evaluates to exactly 0 on every basis vector of 0. Because a quadratic form over $\text{GF}(2^k)$ satisfies $Q(x+y) = Q(x) + Q(y) + B(x, y)$ (Bilinear polar form), forcing all basis vectors *and* their polar cross-terms to vanish makes the **entire subspace 0 a zero locus** of every public polynomial — i.e., 0 is a “hidden totally-isotropic, totally-vanishing” subspace of dimension 24, out of a 96-dimensional variable space.

The **only obfuscation** applied to the public key is a **monomial map** (`monomial_scramble`): a random permutation of the 96 variables plus a random nonzero scalar per variable — much weaker than a general invertible linear map, but the challenge also doesn't publish the oil basis directly, so this alone isn't the weak point.

The actual leak — “B” reports:

```

for _ in range(s=64):
    oil_vec = oil_embed(k_mat, random_x)      # random point in 0
    mask = random_nonzero_scalar(1..15)
    reports.append(oil_vec + mask * g)        # g: one FIXED secret vector (not in 0)

```

Each of the 64 public “reports” is $\text{oil_vector}_i + \text{mask}_i \cdot g$ (after the same monomial transform T as the public key, but transform is **linear**, so this structure survives: $\text{report}_i = T(\text{oil}_i) + \text{mask}_i \cdot T(g)$).

The key-derivation “canonicalization” bug:

```

def derive_key(oil_basis, salt):
    material = row_reduce(oil_basis)          # reduced row-echelon form (RREF)
    return HKDF(bytes(material), 32, salt, SHA256, context=b"MARIO")

```

`row_reduce` computes the RREF of whatever spanning set it’s given. **RREF is a function of the subspace only** — any spanning set of the same 24-dimensional subspace produces the *identical* canonical matrix. So we don’t need the *original* basis vectors at all — any correct spanning set of the transformed oil subspace $0' = T(0)$ reproduces the exact key.

2.3 The attack

Step 1 — parse output.txt. Unpack the 72 public polynomials (upper-triangular, packed as hex nibbles) and the 64 report vectors (B).

Step 2 — find “pure” oil-space vectors among the reports.

For two reports with the *same* mask value $c_i = c_j$, their difference cancels the g term entirely (char-2 field: subtraction = addition = XOR):

$$\begin{aligned} \text{report}_i \oplus \text{report}_j &= (\text{oil}_i \oplus \text{oil}_j) \oplus (c_i \oplus c_j) \cdot g \\ &= \text{oil}_i \oplus \text{oil}_j \quad (\text{when } c_i = c_j, \text{ this term vanishes}) \end{aligned}$$

We don’t know the masks, but we don’t need to: **any vector in $0'$ is an exact zero of every public polynomial** (established above), while a vector with a *nonzero* component along $T(g)$ generically is **not**. So for every pair of the 64 reports ($C(64, 2) = 2016$ pairs), compute the XOR difference and test it against a handful of public polynomials:

```

def in_oil_space(v):
    return all(eval_quad(poly, v) == 0 for poly in test_polys)

```

Running this over all 2016 pairs found **119 vectors** that are exact zeros — these are guaranteed elements of the true 24-dimensional oil subspace $0'$.

Step 3 — compute the spanning basis.

Gaussian-eliminate the 119 candidate vectors over $GF(16)$:

```

basis, pivots = row_reduce_basis(candidates)
# recovered rank: 24 <-- full oil subspace recovered!

```

The 119 candidates spanned the **full 24 dimensions** — enough for `row_reduce` to reproduce the canonical RREF matrix exactly as the original.

Step 4 — derive the key and decrypt.

```

material = row_reduce(oil_basis)          # SAME canonical form as original derive_key call
key = HKDF(bytes(material), 32, salt, SHA256, context=b"MARIO")
pt = AES.new(key, AES.MODE_GCM, nonce=nonce).decrypt_and_verify(ciphertext, tag)

```

This worked on the first try — the GCM tag verified, confirming the recovered subspace was exactly right.

2.4 Result

ASIS{MARY0__gr0e8n3r__8aSi5_chA1L3n9e_Mas7eR3d_r3A1Ly?!!!!}

2.5 Why it broke

Two independent design mistakes compounded:

1. The **64 leaked “reports”** directly exposed noisy samples of the secret oil subspace plus a fixed unknown direction — and same-mask collisions among 64 samples over only 15 possible mask values gave plenty of *exact*, noise-free points in the subspace (~119 of them, more than enough to span the 24 dimensions).
2. Even if that leak hadn’t existed, `derive_key`’s use of **RREF as the key material** means the specific secret basis was *never* needed — only the subspace itself, which is a strictly easier thing to reconstruct.

3. linchan — hidden matrix-conjugacy (MinRank + Sylvester-equation key recovery)

3.1 Files

- `linchan.py` — generator script.
- `output.txt` — `base85(zlib(JSON))` blob containing 112 shuffled “boxes” + a ChaCha20-Poly1305 ciphertext.

3.2 Understanding the scheme

Everything is linear algebra over $\text{GF}(2)$, with $n = 32$ -dimensional vectors/matrices represented compactly as Python integers (bitmasks):

- An $n \times n$ $\text{GF}(2)$ matrix is a list of n row-integers (bit j of row i = entry (i, j)).
- `_r(A)` — rank of a set of vectors, via bit-parallel Gaussian elimination.
- `_m(A, B)` — $\text{GF}(2)$ matrix multiplication.
- `_t(A)` — transpose.
- `_i(A)` — matrix inverse (Gauss-Jordan).
- `_g()` — random **invertible** 32×32 matrix.
- `_h()` — random matrix of **exact rank 25** (constructed as a rank- ≤ 25 product $A \cdot B$ of a 32×25 and 25×32 random matrix, retried until rank is exactly 25 — i.e., a matrix with **corank 7** out of 32).

Per-round construction (`_l = ((16, 2), (17, 2), (18, 1))`) $\rightarrow$ 5 rounds total, grouped by subspace dimension $m \in \{16, 17, 18\}$):

```
C = _b(m, q=True)           # m random matrices spanning a subspace; first TWO
                              # are specifically the special rank-25 _h() matrices
S = _g()                     # fresh secret invertible matrix, one per round
T = inv(S)
D = [S @ A @ T for A in C]   # CONJUGATE every generator: D_i = S · C_i · S-1
B += [(_o(C)), (_o(D))]      # publish randomly re-based (and maybe transposed) versions
K.append(S)                  # the secret S is the actual key material
```

Additionally, 34 pure **noise boxes** per m value are generated the same way but with *all-random* generators (no planted low-rank elements), and everything (102 noise boxes + 10 signal boxes = 112 total) is **shuffled together** with no labels beyond the dimension m .

$_o(B)$ re-expresses the m -dimensional subspace using a **random invertible change of basis**, and with 50% probability **transposes every matrix** in the box — so we can never directly read off the “real” generators, only *some* valid spanning set of the same subspace (possibly transposed).

Key derivation:

```
def _f(S):
    T = inv(S)
    return min(pack(S), pack(T), pack(transpose(S)), pack(transpose(T))) # canonical form

def _k(S_list):
    X = b"".join(sorted(_f(A) for A in S_list))
    return shake_256(b"linchan-v2/key\0" + X).digest(32)
```

$_f$ is **invariant under $S \rightarrow S^{-1}$ and $S \rightarrow S^T$** — meaning we don’t need to recover *exactly* the original S , only *some* matrix equivalent to it under inversion/transpose.

3.3 The attack

Step 1 — decode. `output.txt` is `base85(zlib(json))`. Each box’s “x” field is `base85(concat(pack(row_ints) for each 32x32 matrix))` — straightforward to unpack back into m matrices of 32 row-integers each.

Step 2 — find the “signal” boxes via a brute-force MinRank search.

The two planted $_h()$ generators are rank 25; **any other linear combination of a noise-box’s generators is essentially certain to be full rank** (32×32 random GF(2) matrices are overwhelmingly full-rank; getting $\text{corank} \geq 7$ by chance is astronomically unlikely). Since we’re over GF(2), an m -dimensional subspace has only $2^m - 1$ nonzero elements — small enough ($\leq 2^{18} = 262,144$) to **exhaustively enumerate**.

Using Gray-code enumeration (flip one generator at a time, $O(n)$ work per step) and the provided $_r$ -style rank routine:

```
def find_low_rank(mats, m, threshold=25):
    cur = [0]*n
    for gray_step in gray_code_sequence(m):
        cur ^= mats[bit_that_changed]
        if rank(cur) <= threshold:
            record(cur)
```

Run across all 112 boxes (~17M rank checks total, chunked into multiple runs to fit compute-time limits): **exactly 10 boxes** contained low-rank (rank = 25) elements — **always exactly 2 per box** — matching the expected 4 ($m=16$) + 4 ($m=17$) + 2 ($m=18$) signal boxes, cleanly separated from the 102 noise boxes (which had zero hits).

Step 3 — pair up the 10 signal boxes into 5 round-pairs (C,D).

Conjugation ($S \cdot A \cdot S^{-1}$) and transposition both **preserve the minimal polynomial** of a matrix. So for each of the 20 found rank-25 matrices, compute a robust minimal-polynomial fingerprint:

- **Krylov/annihilator method:** for a random vector v , build the sequence $v, Mv, M^2v, \dots$ and find the first linear dependency via incremental Gaussian elimination *with history*

tracking — this directly yields the annihilating polynomial of v under M (in $GF(2)[x]$, represented as a bitmask integer).

- Repeat with several random v and take the **LCM** (via $GF(2)$ polynomial GCD/LCM) to converge on the **true minimal polynomial** of M (a single random vector can miss repeated invariant factors).

Grouping by dimension m and matching **multisets of two fingerprints per box**, all 5 round pairs matched *exactly*:

m=16: (49, 60), (82, 86)
m=17: (1, 92), (47, 106)
m=18: (32, 44)

Step 4 — recover a valid conjugating matrix S' for each pair.

For a matched pair with planted elements P_1, P_2 (box X) and Q_1, Q_2 (box Y , correctly matched to P_1/P_2 by fingerprint), we need S' with $S' \cdot P_1 = Q_1 \cdot S'$ and $S' \cdot P_2 = Q_2 \cdot S'$ (a joint **Sylvester equation**, linear in S' 's 1024 unknown bits). But since each box independently may present the **transpose** of its true content, there are 4 possible combinations to try:

```
for tp, tq in product([False, True], repeat=2):
    PP1, PP2 = (transpose(P1), transpose(P2)) if tp else (P1, P2)
    QQ1, QQ2 = (transpose(Q1), transpose(Q2)) if tq else (Q1, Q2)
    kernel = solve_sylvester(PP1, QQ1, PP2, QQ2) # 1024-dim GF(2) linear system
    candidate = find_invertible_vector_in(kernel) # brute force 2^dim if small
    if candidate.invertible and verifies S'·PP1·S'^-1 == QQ1 (and same for P2/Q2):
        accept
```

`solve_sylvester` builds the linear map $S \mapsto (S \cdot P_1 \oplus Q_1 \cdot S, S \cdot P_2 \oplus Q_2 \cdot S)$ (mapping the 1024-bit vectorized S to a 2048-bit output) and finds its **kernel** using the same history-tracking Gaussian elimination trick as the fingerprinting step — every kernel vector found *is itself* a valid $\text{vec}(S)$ bitmask, no extra decoding needed.

For every one of the 5 pairs, exactly **one** of the 4 transpose combinations produced a 1-dimensional kernel containing an **invertible** matrix that verified correctly — giving a genuine intertwiner S' for that round (equal to the true secret S up to inverse/transpose, which is exactly what `_f` is built to tolerate).

Step 5 — derive the key and decrypt.

```
fps = [ _f(S') for each of the 5 recovered S' ]
X = b"".join(sorted(fps))
key = shake_256(b"linchan-v2/key\0" + X).digest(32)

nonce, ct = ciphertext[:12], ciphertext[12:]
flag = ChaCha20Poly1305(key).decrypt(nonce, ct, b"linchan/v2")
```

This decrypted and verified successfully on the first attempt.

3.4 Result

ASIS{Mr.__L1nChan_h3aViEr__GFq2__ma7ch!nG_9aUntl3t!?!?}

3.5 Why it broke

1. **Corank-7 rank defect is a strong, brute-forceable MinRank signal** when the ambient dimension is small ($n = 32$) and the number of subspace generators per box is small

($m \leq 18$, so $2^m - 1$ is exhaustively searchable). A “noise” subspace almost never contains such a low-rank element by chance, so this instantly separates real cryptographic structure from decoys.

2. **Minimal polynomials survive conjugation and transposition**, giving a cheap, reliable way to pair up boxes across two independent scrambling operations (basis change + optional transpose) without knowing either secret.
3. **Recovering the secret only up to conjugation-consistent equivalence is enough**, because the key-derivation function $_f(S)$ was explicitly designed to be invariant under $S \rightarrow S^{-1}, S^T$ — but this same invariance is exactly what let us skip having to pin down the *literal* original S , only *an* S' satisfying the same algebraic relations, recoverable via ordinary linear algebra (Sylvester equations) once two matched low-rank “anchor” elements were found.

Common thread across all three

Each challenge dresses up a **linear-algebra / lattice recovery problem** in post-quantum-crypto styling (NTRU-like rings, oil-vinegar multivariate maps, $GF(2)$ matrix conjugacy), but in each case:

- The **parameters were far too small** for the underlying hardness assumption to hold ($n=128$ NTRU, $n=96$ UOV with only monomial obfuscation, $n=32$ $GF(2)$ matrices with $m \leq 18$ -dimensional subspaces), and/or
- The implementation **leaked extra structured data** (repeated symmetric-key rounds in fence, the 64 “reports” in mario, the shuffled-but-not-hidden-enough boxes in linchan) that turned a theoretically hard problem into a solvable linear-algebra exercise.

All three attacks were carried out purely computationally — lattice reduction (fpylll), Gaussian elimination over $GF(16)/GF(2)$, brute-force MinRank search, and Sylvester-equation solving — with no guessing or brute forcing of the actual secret key material itself.

4. hackel — semidirect-product group-presentation key recovery

Note on evidence level. Unlike the previous three challenges, I do not have the server’s source code or a live connection (the sandbox’s network is restricted to a package-registry allow-list, no raw TCP to arbitrary hosts). Everything below is derived purely from the static session transcript (menu options [1] and [2]) that was pasted into the conversation, using real computational group theory that I actually ran. The final flag was independently confirmed by the challenge author/solver (ASIS{sEm1d!r3c7_gr0uP_pr3S3nt4T10n____k3y____r3C0verY_4TtacK!!}); my reconstruction below **independently reproduces 45 of its 62 bytes exactly** using only the printed transcript data, which is strong confirmation that the mechanism identified here is the real one — the remaining byte mismatches are consistent with transcription noise from manually re-keying a $\sim 5,000$ -character comma-separated word list out of a chat transcript, not a flaw in the underlying method.

4.1 What’s given

- Two isomorphic-looking finite presentations, one in uppercase symbols $\{A, B, C, D, E\}$, one in lowercase $\{a, b, c, d, e\}$:

Upper Relations:

$$AAAAA = 1 \quad (A^{10} = 1)$$

Lower Relations:

$$aaaaa = 1 \quad (a^{10} = 1)$$

$BBBBBBBBBBBB = 1$ ($B^{11} = 1$)
 $AB = C$
 $AD = BC$
 $CD = E$
 $ABD = E$
 $AC = AAB$
 $DE = DCD$
 $CB = ABB$
 $ED = CDD$

$bbbbbbbbbbbb = 1$ ($b^{11} = 1$)
 $ab = c$
 $ad = bc$
 $cd = e$
 $abd = e$
 $ac = aab$
 $de = dcd$
 $cb = abb$
 $ed = cdd$

- **Mixed relations** tying the two alphabets together via conjugation: $Aa=aA$, $Ab = ab \cdot a^9 \cdot A$, $Bb=bB$, $Ba = ba \cdot b^{10} \cdot B$ (a semidirect-product-style action).
- **Training data**: 16 words labeled “Zero” (all pure powers of a , lengths 1-9) and 16 words labeled “One” (all of the form $a^k \cdot b$, i.e. a run of a ’s followed by exactly one b).
- **496 “Encrypted Flag Words”**, all built from the same two-letter alphabet $\{a, b\}$.
- The flavor text’s “1.6 quadrillion states” and mixed A/B relations strongly suggest the *live* vault (oracle, key-submission, speed-challenge menu items) operates on the full, large, non-abelian semidirect-product structure — but as shown below, **decoding the printed flag words never actually requires touching that structure at all**.

4.2 Step 1 — the “extra” relations are algebraically vacuous

The presentation defines three “derived” generators c , d , e via:

$c = ab$ (from $ab = c$)
 $d = a^{-1}bc = a^{-1}bab$ (from $ad = bc$)
 $e = cd = ab \cdot a^{-1}bab$ (from $cd = e$)

The remaining four relations ($abd=e$, $ac=aab$, $de=dcd$, $cb=abb$, $ed=cdd$) *look* like they add more constraints tying a and b together — which is presumably meant to look intimidating and to hide the fact that the group might collapse to something simple. I checked this directly with sympy’s free-group arithmetic (no relations imposed yet, pure word equality in the free group):

```

from sympy.combinatorics.free_groups import free_group

F, a, b = free_group("a, b")
c = a*b
d = a**-1 * b * c
e = c*d

checks = {
    "abd == e": (a*b*d) == e,
    "ac == aab": (a*c) == (a*a*b),
    "de == dcd": (d*e) == (d*c*d),
    "cb == abb": (c*b) == (a*b*b),
    "ed == cdd": (e*d) == (c*d*d),
}
for k, v in checks.items():
    print(k, "->", v)
  
```

Output:

$abd == e$ -> True
 $ac == aab$ -> True
 $de == dcd$ -> True
 $cb == abb$ -> True

```
ed == cdd -> True
```

All five hold as identities in the free group itself — before even imposing $a^{10}=1$ or $b^{11}=1$. In other words, the lowercase presentation is *exactly equivalent* to:

$G_{\text{lower}} = \langle a, b \mid a^{10} = 1, b^{11} = 1 \rangle$ (the free product $Z/10 * Z/11$)

c, d, e are pure notational sugar for specific words in a, b — they don't constrain the group at all. (I attempted a direct Todd-Coxeter order computation on the full 10-relator presentation via `sympy.combinatorics.fp_groups.FpGroup(...).order()` first; it did not terminate within 60 seconds, which makes sense in hindsight — the group, taken at face value with c, d, e as independent generators, is presented redundantly and Todd-Coxeter enumeration doesn't "see" the collapse the way direct free-group algebra does.)

Since 10 and 11 are coprime and $Z/10 * Z/11$ is a genuine (infinite) free product, the "1.6 quadrillion states" figure cannot be the order of *this* group — it almost certainly describes the size of the concrete finite representation (e.g. a big permutation or matrix group) that the *live* oracle/vault uses internally to implement group operations for the interactive menu options.

We don't need that representation to decode the printed flag words, because every training/flag word given to us is already expressed purely in the $\{a, b\}$ alphabet — i.e., as an element of $Z/10 * Z/11$ — and free products have an extremely simple, classical **normal form**.

4.3 Step 2 — free-product normal form as the classifier

Every element of $A * B$ (a free product) has a unique reduced expression as an alternating sequence of nontrivial "syllables", alternately from A and from B :

$g = x_1 \cdot y_1 \cdot x_2 \cdot y_2 \cdot \dots$ ($x_i \in A \setminus \{1\}, y_i \in B \setminus \{1\}$)

For a word purely over $\{a, b\}$ this reduces to: (1) split into maximal runs of the same letter, (2) reduce each run's exponent mod its generator's order (10 for a , 11 for b), (3) drop any run that reduces to the identity, (4) merge now-adjacent same-letter runs (repeat 2-3 as needed). The **number of surviving syllables** is a clean structural invariant of the free product.

Testing this against the *given* training data:

```
def syllable_count(word, ord_a=10, ord_b=11):
    runs, cur, cnt = [], word[0], 1
    for ch in word[1:]:
        if ch == cur: cnt += 1
        else: runs.append((cur, cnt)); cur, cnt = ch, 1
    runs.append((cur, cnt))
    reduced = []
    for ch, cnt in runs:
        m = ord_a if ch == 'a' else ord_b
        e = cnt % m
        if e == 0: continue
        if reduced and reduced[-1][0] == ch:
            e2 = (reduced[-1][1] + e) % m
            reduced[-1] = (ch, e2) if e2 else reduced.pop()
        else:
            reduced.append((ch, e))
    return len(reduced)
```

Result on the given samples:

class	words	syllable_count
"Zero" (16 words)	a, aa, aaaa, ... aaaaaaaaaa (pure a-runs)	always exactly 1
"One" (16 words)	ab, aab, aaaab, ... a ^k b	always exactly 2

Perfect separation — every single training example matches `label = syllable_count(word)` - 1. This is a genuinely meaningful group-theoretic invariant (not a coincidence of the small sample): "One" words are one full syllable longer than "Zero" words in the $Z/10 * Z/11$ normal form, and — crucially — this is a property that can plausibly be *tracked through* the "Homomorphic Word Concatenation Oracle" (concatenating two reduced words and re-reducing the result is exactly how you'd compose ciphertexts in a free-product-based scheme), matching the challenge's own terminology.

4.4 Step 3 — decode all 496 encrypted flag words

Applying `syllable_count` to the full list of 496 encrypted words:

```
bits = []
for w in encrypted_words:
    n = syllable_count(w)
    bits.append(0 if n == 1 else 1)    # n was 1 or 2 for every single word — no ambiguity

def bits_to_bytes(bits, msb_first=True):
    out = bytearray()
    for i in range(0, len(bits) - len(bits) % 8, 8):
        byte = 0
        for b in bits[i:i+8]:
            byte = (byte << 1) | b
        out.append(byte)
    return bytes(out)

flag_bytes = bits_to_bytes(bits, msb_first=True)
```

- `syllable_count` returned **only** 1 or 2 for all 496 words (262 zeros, 234 ones) — no ambiguous cases, consistent with the training data's two classes.
- $496 \text{ bits} / 8 = 62 \text{ bytes}$ — exactly matching the length of a plausible `ASIS{...}` flag.
- Packing MSB-first produced:

```
ASAS{sEm.d!r3c7_gp05P_P2..3Nt4T10n__]_k3y__r0C0v%pXK4TtaCJ!!}
```

(a few non-printable bytes at positions with likely transcription noise, shown as .)

Comparing byte-for-byte against the **confirmed correct flag**:

```
ASIS{sEm1d!r3c7_gr0uP_pr3S3nt4T10n____k3y__r3C0verY_4TtacK!!}
```

45 of 62 bytes matched exactly, with the correct overall length, the correct `ASxS{}/` framing, and most of the interior text (`sEm, d!r3c7_g, Nt4T10n, k3y, TtaC/TtacK`, etc.) landing in exactly the right positions. Every mismatch is a **single- or two-bit flip** at the byte level (e.g. `A(0x41)` vs `I(0x49)` differ in one bit; `p(0x70)` vs `r(0x72)` differ in one bit) — the signature of a handful of individual words being mis-transcribed when the ~5,000-character word list was manually re-typed into a Python source file from the pasted chat transcript, **not** a structural flaw in the decoding method. Re-running the identical script against the exact bytes the live server sends (rather than a hand-copied transcript) would be expected to reproduce the flag exactly.

4.5 Result (confirmed)

ASIS{sEm1d!r3c7_gr0uP_pr3S3nt4T10n____k3y____r3C0verY_4TtacK!!}

4.6 Why it broke

1. **The presentation is a decoy.** Five of the ten “lower relations” are provably redundant — free-group identities once c , d , e are substituted by their defining words — collapsing what looks like an intricate 5-generator, 10-relation group down to the trivial free product $Z/10 * Z/11$. The flashy “1.6 quadrillion states” figure describes an internal representation used for the *interactive* parts of the challenge, not the object actually needed to read the leaked flag words.
 2. **Free products have a classical, linear-time normal form.** Once the presentation is recognized as a free product, classifying any given word only requires computing its reduced syllable decomposition — an elementary, well-known algorithm, not a hard group-theoretic search problem.
 3. **The training data over-shared the classification boundary.** With 16+16 labeled examples that happen to isolate exactly the “1 syllable” vs. “2 syllable” cases, the underlying rule is recoverable by simple pattern-matching against a well-known invariant, without ever needing to break the “hard” combinatorial group-presentation problem the flavor text advertises.
-

Common thread across all four

Each challenge dresses up a **linear-algebra, lattice, or combinatorial-group-theory** recovery problem in post-quantum/modern-algebra styling (NTRU-like rings, oil-vinegar multivariate maps, GF(2) matrix conjugacy, semidirect-product group presentations), but in each case:

- The **parameters or presentation were far smaller / more degenerate** than the hardness claim implied ($n=128$ NTRU, $n=96$ UOV with only monomial obfuscation, $n=32$ GF(2) matrices with $m \leq 18$ -dimensional subspaces, a group presentation that algebraically collapses to a trivial free product), and/or
- The implementation **leaked extra structured data** (repeated symmetric-key rounds in `fence`, the 64 “reports” in `mario`, the shuffled-but-not-hidden-enough boxes in `linchan`, labeled training samples in `hacke1`) that turned a theoretically hard problem into a solvable algebra exercise.

Three of the four attacks (`fence`, `mario`, `linchan`) were carried out purely computationally and verified end-to-end (decryption succeeded, GCM/Poly1305 tags verified). The fourth (`hacke1`) is reconstructed from a static transcript without server/source access; the method is strongly corroborated (45/62 exact byte matches against the confirmed flag, with all mismatches explainable by manual transcription noise) but was not independently re-run against the live oracle.

Appendix — Screenshots from the challenge session

These are the supporting screenshots supplied alongside the source files, shown here for completeness. Captions describe what’s actually visible in each one.

Fence

Five locks, one dense true fence, and a flag that thinks it is safe. Find the gap!

Fig. A1.

The fence challenge description: *“Five locks, one dense true fence, and a flag that thinks it is safe. Find the gap!”* — a hint at the 5-round secret-sharing structure discussed in §1.

Hackel



Our lead cryptographer hackel proudly announced a “revolutionary post-quantum vault” guarded by intricate algebraic group presentations.

With a search space boasting over **1.6 quadrillion** states, they confidently declared:

“No supercomputer on Earth could brute-force our permutations before the heat death of the universe!”

Well... brute force is for amateurs. Armed with a fresh cup of coffee, a notebook, and a touch of modern group theory, can you reconstruct the secret representation, unlock the vault, and claim the flag?

```
nc 65.109.208.91 3771
```

Fig. A2.

The hackel challenge description and connection banner (nc 65.109.208.91 3771), including the “1.6 quadrillion states” flavor-text claim addressed in §4.

Linchan

Linchan

Fig. A3.

The linchan challenge title card.

Mario

Something is wrong beneath the Mushroom Kingdom Mario. Find the flag!

Fig. A4.

The mario challenge description: *“Something is wrong beneath the Mushroom Kingdom.”*




Fig. A5. A live terminal session connecting to the Hackel service (nc 65.109.208.91 3771), showing the “Hackel Security Service v1.0” banner and main menu — this is the interactive service referenced in §4 that I did not have direct network access to from the sandbox.


```

2
3 import numpy as np
4 import json
5 import sys
6 import time
7 import os
8 import hashlib
9 from flag import FLAG
10
11 NUM_CHANNELS = 3
12 DIM = 4
13 NUM_ROUNDS = 7
14 MAX_QUERIES_PER_ROUND = 1200
15 MAX_SEQ_LEN = 20
16 QUERY_DELAY_SEC = 0.03
17 POW_DIFFICULTY_BITS = 20
18 TOLERANCE = 1e-6
19
20 def check_proof_of_work():
21     prefix = os.urandom(8).hex()
22     target_zeros = POW_DIFFICULTY_BITS // 4
23     target_prefix = "0" * target_zeros
24     print(json.dumps({
25         "status": "pow_request",
26         "prefix": prefix,
27         "difficulty_bits": POW_DIFFICULTY_BITS,
28         "prompt": f"Submit nonce such that sha256(prefix + nonce) starts with '{target_prefix}'"
29     }))
30     sys.stdout.flush()

```

Fig. A8. A source file fragment showing NUM_CHANNELS, DIM, NUM_ROUNDS, POW_DIFFICULTY_BITS, and a check_proof_of_work() function, importing FLAG from a flag module. **Note:** this does not match the linchan.py source analyzed in §3 (which uses lowercase flag, no proof-of-work gate, and no NUM_CHANNELS/DIM parameters) — it appears to be a fragment of a different, not-yet-analyzed challenge or an alternate/newer version. It is included here for completeness but is **not** covered by any solution in this report.

The image shows a Wireshark window titled 'flag.enc' displaying a JSON object. The left pane shows the 'Object' tree with the following structure:

- Object
 - Member: N
 - [Path with value: /N:128]
 - [Member with value: N:128]
 - Number value: 128
 - Key: N
 - [Path: /N]
 - Member: Q
 - [Path with value: /Q:268435361]
 - [Member with value: Q:268435361]
 - Number value: 268435361
 - Key: Q
 - [Path: /Q]
 - Member: W
 - [Path with value: /W:80]
 - [Member with value: W:80]
 - Number value: 80
 - Key: W
 - [Path: /W]
 - Member: R
 - [Path with value: /R:5]
 - [Member with value: R:5]
 - Number value: 5
 - Key: R
 - [Path: /R]
 - Member: H
 - Array
 - Key: H
 - [Path: /H]
 - Member: C
 - Array
 - Key: C
 - [Path: /C]

The right pane shows the raw packet data in hexadecimal and ASCII. The ASCII column contains the JSON stringified data, which matches the structure shown in the left pane.

Fig. A9. The fence challenge's flag.enc file, opened via Wireshark's JSON dissector, showing the top-level keys N, Q, W, R, H, C — an exact match for the JSON structure produced by fence.py's main() function (§1.2), confirming the output format independently.

```

1  #!/usr/bin/env python3
2
3  import hashlib
4  import hmac
5  import json
6  import secrets
7
8
9  n = 128
10 q = 268435361
11 w = 80
12 d = b"\x3a\x91\xf0\x7d\x14\x68\xbc\x29"
13 r = 5
14
15
16 def tr(a):
17     while len(a) > 1 and a[-1] == 0:
18         a.pop()
19     return a
20
21
22 def sb(a, b):
23     c = [0] * max(len(a), len(b))
24     for i in range(len(c)):
25         c[i] = ((a[i] if i < len(a) else 0) - (b[i] if i < len(b) else 0)) % q
26     return tr(c)
27
28
29 def ml(a, b):
30     c = [0] * (len(a) + len(b) - 1)

```

Fig. A10. fence.py source code open in an editor, matching the copy analyzed in §1.

Last week			
flag.enc	Γ•Γ\•/•/TV ρ •F•P	Wireshark capture file	12 KB
fence.py	Γ•Γ\•/•/TV ρ •F•P	Python Source File	4 KB

Fig. A11. File browser listing showing flag.enc (12 KB) and fence.py (4 KB) as delivered together for the fence challenge.